

# Assignment-1-Part-2-1

April 25, 2025

## 1 Assignment 1 Part 2.1

### 1.1 Import necessary modules

```
[1]: import torch
import torch.nn as nn
import os
import mlflow
from datetime import datetime
import matplotlib.pyplot as plt
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
```

### 1.2 Layer class

```
[2]: class Layer:
    def __init__(self, neurons_size: int, inputs_size: int, activation: str = 'none',
        ↪ 'none',
        leaky_slope: float = 0.01, softmax_dim: int = 0):
        """
        Initialize the layer with random weights and biases

        Parameters:
            neurons_size: int
                The number of neurons in the layer
            inputs_size: int
                The number of inputs to the layer
            activation: str
                The activation function to use
            leaky_slope: float
                The slope of the leaky relu activation function
            softmax_dim: int
                The dimension of the softmax activation function
        """

        # initialize weights and biases with random values
```

```

        # use the He initialization method to initialize the weights if the
        ↪ activation function is relu or leaky relu else use the Xavier initialization
        ↪ method
        if activation == 'relu' or activation == 'leaky_relu':
            self._weights: torch.Tensor = torch.randn(
                neurons_size, inputs_size) * torch.sqrt(2 / torch.
            ↪ tensor(inputs_size, dtype=torch.float32))
        else:
            self._weights: torch.Tensor = torch.randn(
                neurons_size, inputs_size) * torch.sqrt(1 / torch.
            ↪ tensor(inputs_size, dtype=torch.float32))
            self._biases: torch.Tensor = torch.zeros(neurons_size)
            self._activation: str = activation
            self._leaky_slope: float = leaky_slope
            self._softmax_dim: int = softmax_dim

            self._activation_function: dict[str, callable] = {
                'sigmoid': nn.Sigmoid(),
                'tanh': nn.Tanh(),
                'relu': nn.ReLU(),
                'leaky_relu': nn.LeakyReLU(negative_slope=self._leaky_slope),
                'softmax': nn.Softmax(dim=self._softmax_dim),
                'none': nn.Identity()
            }

    def set_weights(self, weights: torch.Tensor):
        """
        Set the weights of the layer

        Parameters:
            weights: torch.Tensor
                The weights to set
        """
        self._weights = weights

    def get_weights(self) -> torch.Tensor:
        """
        Get the weights of the layer

        Returns:
            torch.Tensor
                The weights of the layer
        """
        return self._weights

    def set_biases(self, biases: torch.Tensor):
        """

```

```

        Set the biases of the layer

Parameters:
    biases: torch.Tensor
        The biases to set
    """
    self._biases = biases

def get_biases(self) -> torch.Tensor:
    """
    Get the biases of the layer

Returns:
    torch.Tensor
        The biases of the layer
    """
    return self._biases

def forward(self, inputs: torch.Tensor) -> torch.Tensor:
    """
    Forward pass

Parameters:
    inputs: torch.Tensor
        The inputs to the layer
    activation: str
        The activation function to use
Returns:
    torch.Tensor
        The outputs of the layer
    """
    # calculate the sum of the inputs multiplied by the weights and add the
    ↪biases
    if inputs.dim() == 2:
        sum: torch.Tensor = torch.matmul(
            inputs, self._weights.t()) + self._biases
    else:
        sum: torch.Tensor = torch.matmul(
            inputs, self._weights) + self._biases

    if self._activation in self._activation_function:
        return self._activation_function[self._activation](sum)
    else:
        raise ValueError(
            f"Activation function {self._activation} not found")

```

### 1.3 Custom Layer class

```
[3]: class CustomLayer(nn.Module):
    def __init__(self, neurons_size: int, inputs_size: int, activation: str = 'none',
        ↪leaky_slope: float = 0.01, softmax_dim: int = 0):
        super(CustomLayer, self).__init__()

        '''
        Initialize the custom layer with the given neurons size, inputs size,
        activation, leaky slope, and softmax dimension
        '''

        # initialize the custom layer
        self.layer = Layer(neurons_size, inputs_size,
            activation, leaky_slope, softmax_dim)

        # register the weights and biases as parameters to PyTorch
        self.weights = nn.Parameter(self.layer.get_weights())
        self.biases = nn.Parameter(self.layer.get_biases())

        # point the layer parameters to the PyTorch parameters
        self.layer.set_weights(self.weights)
        self.layer.set_biases(self.biases)

        # store the activation function parameters
        self.activation = activation
        self.leaky_slope = leaky_slope
        self.softmax_dim = softmax_dim

    def forward(self, inputs: torch.Tensor) -> torch.Tensor:

        # use custom layer forward method
        return self.layer.forward(inputs)
```

### 1.4 Perceptron class

```
[4]: class Perceptron(nn.Module):
    def __init__(self, input_size: int, hidden_size: int, output_size: int,
        ↪hidden_activation: str = 'relu',
        ↪output_activation: str = 'none', leaky_slope: float = 0.01,
        ↪softmax_dim: int = 0):
        super(Perceptron, self).__init__()

        '''
```

```

        Initialize the perceptron with the given input size, hidden size,
        ↪output size,
        hidden activation, output activation, leaky slope, and softmax dimension

Parameters:
    input_size: int
        The size of the input layer
    hidden_size: int
        The size of the hidden layer
    output_size: int
        The size of the output layer
    hidden_activation: str
        The activation function to use for the hidden layer
    output_activation: str
        The activation function to use for the output layer
    leaky_slope: float
        The slope of the leaky relu activation function
    softmax_dim: int
        The dimension of the softmax activation function
'''

# create a list of layer sizes including the input and output sizes
layer_sizes: list[int] = [input_size] + hidden_size + [output_size]
layers: list[CustomLayer] = []

# create hidden layers
for i in range(len(layer_sizes) - 2):
    layers.append(CustomLayer(
        layer_sizes[i+1], layer_sizes[i], hidden_activation,
        ↪leaky_slope, softmax_dim))

# create output layer
layers.append(CustomLayer(
    layer_sizes[-1], layer_sizes[-2], output_activation, leaky_slope,
    ↪softmax_dim))

# store the layers
self.layers = nn.ModuleList(layers)

def forward(self, inputs: torch.Tensor) -> torch.Tensor:
    '''
    Forward pass

Parameters:
    inputs: torch.Tensor
        The inputs to the perceptron

```

```

Returns:
    torch.Tensor
    The outputs of the perceptron
'''

# process the input through the layers
x = inputs
for layer in self.layers:
    x = layer(x)
return x

```

## 1.5 Trainer class

```

[5]: class Trainer:
    def __init__(self,
                model,
                criterion,
                optimizer,
                train_loader,
                val_loader=None,
                test_loader=None,
                device=None,
                scheduler=None,
                checkpoint_dir='./checkpoints',
                experiment_name=None):
        """
        Initialize the trainer with model, criterion, optimizer, and data_
        ↪ loaders

        Parameters:
            model: nn.Module
                The model to train
            criterion: loss function
                The loss function to use
            optimizer: torch.optim
                The optimizer to use
            train_loader: DataLoader
                The data loader for training data
            val_loader: DataLoader
                The data loader for validation data
            test_loader: DataLoader
                The data loader for test data
            device: torch.device
                The device to use for training
            scheduler: torch.optim.lr_scheduler
                Learning rate scheduler (optional)
            checkpoint_dir: str

```

```

        Directory to save checkpoints
        experiment_name: str
        Name of the experiment for MLflow tracking
    """
    self.model = model
    self.criterion = criterion
    self.optimizer = optimizer
    self.train_loader = train_loader
    self.val_loader = val_loader if val_loader is not None else test_loader
    self.test_loader = test_loader
    self.device = device if device is not None else torch.device(
        'cuda' if torch.cuda.is_available() else 'cpu')
    self.scheduler = scheduler
    self.checkpoint_dir = checkpoint_dir
    self.experiment_name = experiment_name
    self.input_size = None # will be set during first forward pass

    # create checkpoint directory if it doesn't exist
    os.makedirs(checkpoint_dir, exist_ok=True)

    # move model to device
    self.model.to(self.device)

    # initialize training metrics
    self.best_accuracy = 0.0
    self.best_loss = float('inf')
    self.history = {
        'train_loss': [],
        'val_loss': [],
        'val_accuracy': [],
        'learning_rates': []
    }

def train_epoch(self, epoch, num_epochs):
    """
    Train the model for one epoch
    """
    self.model.train()
    running_loss = 0.0

    for i, (inputs, labels) in enumerate(self.train_loader):
        # get input size from first batch if not set
        if self.input_size is None and hasattr(inputs, 'view'):
            self.input_size = inputs.view(inputs.size(0), -1).size(1)

        # prepare inputs
        if hasattr(inputs, 'view'):

```

```

        # for image data, flatten if needed
        inputs = inputs.view(inputs.size(0), -1).to(self.device)
    else:
        inputs = inputs.to(self.device)

    labels = labels.to(self.device)

    # forward pass
    outputs = self.model(inputs)
    loss = self.criterion(outputs, labels)

    # backward and optimize
    self.optimizer.zero_grad()
    loss.backward()
    self.optimizer.step()

    # update statistics
    running_loss += loss.detach().item()

    # print progress
    if (i+1) % 100 == 0:
        print(
            f'Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/{len(self.
↪train_loader)}], Loss: {loss.detach().item():.4f}')

    # calculate average loss for the epoch
    epoch_loss = running_loss / len(self.train_loader)
    self.history['train_loss'].append(epoch_loss)

    # track learning rate
    current_lr = self.optimizer.param_groups[0]['lr']
    self.history['learning_rates'].append(current_lr)

    return epoch_loss

def validate(self, epoch=None):
    """
    Validate the model on validation data
    """
    if self.val_loader is None:
        return None, None

    self.model.eval()
    running_loss = 0.0
    correct = 0
    total = 0

```



```

with torch.no_grad():
    for inputs, labels in self.val_loader:
        # prepare inputs
        if hasattr(inputs, 'view'):
            # for image data
            inputs = inputs.view(inputs.size(0), -1).to(self.device)
        else:
            inputs = inputs.to(self.device)

        labels = labels.to(self.device)

        # forward pass
        outputs = self.model(inputs)
        loss = self.criterion(outputs, labels)

        # update statistics
        running_loss += loss.item()

        # calculate accuracy
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    # calculate average loss and accuracy
    val_loss = running_loss / len(self.val_loader)
    val_accuracy = 100 * correct / total

    if epoch is not None:
        self.history['val_loss'].append(val_loss)
        self.history['val_accuracy'].append(val_accuracy)

    return val_loss, val_accuracy

def test(self):
    """
    Test the model on test data
    """
    if self.test_loader is None:
        return None, None

    return self.validate() # reuse validation code for testing

def save_checkpoint(self, epoch, train_loss, val_loss=None,
    ↪ val_accuracy=None, is_best=False):
    """
    Save a checkpoint of the model
    """

```

```

checkpoint = {
    'epoch': epoch + 1,
    'model_state_dict': self.model.state_dict(),
    'optimizer_state_dict': self.optimizer.state_dict(),
    'loss': train_loss
}

if val_loss is not None:
    checkpoint['val_loss'] = val_loss

if val_accuracy is not None:
    checkpoint['val_accuracy'] = val_accuracy

if self.scheduler is not None:
    checkpoint['scheduler_state_dict'] = self.scheduler.state_dict()

# save regular checkpoint
checkpoint_path = os.path.join(
    self.checkpoint_dir, f'checkpoint_epoch_{epoch+1}.pt')
torch.save(checkpoint, checkpoint_path)

# save as best model if applicable
if is_best:
    best_model_path = os.path.join(
        self.checkpoint_dir, 'best_model.pt')
    torch.save(checkpoint, best_model_path)
    print(f'New best model saved with accuracy: {val_accuracy:.2f}%')

return checkpoint_path

def load_checkpoint(self, checkpoint_path):
    """
    Load a checkpoint
    """
    print(f"Loading checkpoint from {checkpoint_path}")
    checkpoint = torch.load(checkpoint_path, map_location=self.device)

    self.model.load_state_dict(checkpoint['model_state_dict'])
    self.optimizer.load_state_dict(checkpoint['optimizer_state_dict'])

    if 'scheduler_state_dict' in checkpoint and self.scheduler is not None:
        self.scheduler.load_state_dict(checkpoint['scheduler_state_dict'])

    start_epoch = checkpoint['epoch']
    loss = checkpoint.get('loss', 0)
    val_loss = checkpoint.get('val_loss', 0)
    val_accuracy = checkpoint.get('val_accuracy', 0)

```

```

print(
    f"Checkpoint loaded - Epoch: {start_epoch}, Loss: {loss:.4f}, Val_
↳ Loss: {val_loss:.4f}, Val Accuracy: {val_accuracy:.2f}%"

    return start_epoch

def train(self, num_epochs, start_epoch=0, log_to_mlflow=True,
↳ early_stopping_patience=None):
    """
    Train the model for multiple epochs

    Parameters:
        num_epochs: int
            Number of epochs to train for
        start_epoch: int
            Starting epoch (useful when resuming training)
        log_to_mlflow: bool
            Whether to log metrics to MLflow
        early_stopping_patience: int
            Number of epochs to wait for improvement before stopping
    """
    # set mlflow experiment if provided
    if log_to_mlflow and self.experiment_name:
        mlflow.set_experiment(self.experiment_name)

    # initialize early stopping variables
    if early_stopping_patience is not None:
        early_stopping_counter = 0
        best_val_loss = float('inf')

    # start mlflow run if applicable
    run_context = mlflow.start_run(
        run_name=f"{self.model.__class__.__name__}_{datetime.now().
↳ strftime('%Y%m%d_%H%M%S')}") if log_to_mlflow else DummyContextManager()

    # store previous learning rate to detect changes
    prev_lr = self.optimizer.param_groups[0]['lr']

    with run_context:
        # log parameters if using mlflow
        if log_to_mlflow:
            params = {
                "optimizer": self.optimizer.__class__.__name__,
                "learning_rate": self.optimizer.param_groups[0]['lr'],
                "num_epochs": num_epochs
            }

```

```

        # add model architecture params if available
        if hasattr(self.model, 'input_size'):
            params["input_size"] = self.model.input_size
        mlflow.log_params(params)

    print("Starting training...")
    for epoch in range(start_epoch, num_epochs):
        # train for one epoch
        print('')
        train_loss = self.train_epoch(epoch, num_epochs)
        print(
            f'Epoch [{epoch+1}/{num_epochs}], Loss: {train_loss:.4f}')

        # validate
        val_loss, val_accuracy = self.validate(epoch)

        # early stopping check for overfitting
        if early_stopping_patience is not None and val_loss is not None:
            if val_loss < best_val_loss:
                best_val_loss = val_loss
                early_stopping_counter = 0
            else:
                early_stopping_counter += 1
                print(
                    f'EarlyStopping counter: {early_stopping_counter}␣
↳out of {early_stopping_patience}')

                if early_stopping_counter >= early_stopping_patience:
                    print(
                        f'Early stopping triggered after epoch␣
↳{epoch+1}')
                    break

        # log metrics
        if log_to_mlflow:
            mlflow.log_metric("train_loss", train_loss, step=epoch)
            if val_loss is not None:
                mlflow.log_metric("val_loss", val_loss, step=epoch)
            if val_accuracy is not None:
                mlflow.log_metric(
                    "val_accuracy", val_accuracy, step=epoch)

        # print validation results
        if val_loss is not None and val_accuracy is not None:
            print(
                f'Validation - Epoch [{epoch+1}/{num_epochs}], Loss:␣
↳{val_loss:.4f}, Accuracy: {val_accuracy:.2f}%')

```

```

        # check if this is the best model so far
        is_best = False
        if val_accuracy is not None and val_accuracy > self.
↪best_accuracy:
            self.best_accuracy = val_accuracy
            is_best = True

        # save checkpoint
        self.save_checkpoint(
            epoch, train_loss, val_loss, val_accuracy, is_best)

        # update learning rate if scheduler is provided
        if self.scheduler is not None:
            if isinstance(self.scheduler, torch.optim.lr_scheduler.
↪ReduceLROnPlateau):
                self.scheduler.step(val_loss)
            else:
                self.scheduler.step()

        # check if learning rate changed
        current_lr = self.optimizer.param_groups[0]['lr']
        if current_lr != prev_lr:
            print(
                f"Learning rate changed from {prev_lr:.6f} to_
↪{current_lr:.6f}")
            prev_lr = current_lr

    print("Training complete!")

    # final evaluation
    print("\nFinal Evaluation:")
    test_loss, test_accuracy = self.test()
    if test_loss is not None and test_accuracy is not None:
        print(f'Final Test Loss: {test_loss:.4f}')
        print(f'Final Test Accuracy: {test_accuracy:.2f}%')

    # log final metrics
    if log_to_mlflow:
        mlflow.log_metric("test_loss", test_loss)
        mlflow.log_metric("test_accuracy", test_accuracy)

    # save the final model
    model_timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
    final_model_path = os.path.join(
        self.checkpoint_dir, f"final_model_{model_timestamp}.pt")
    torch.save(self.model.state_dict(), final_model_path)

```

```

print(f"Final model saved as {final_model_path}\n")

# log the model with mlflow
if log_to_mlflow and self.input_size is not None:

    model_to_log = self.model.to('cpu')

    # create an input example
    sample_input = torch.rand(1, self.input_size).numpy()

    # create custom pip requirements
    pip_requirements = [
        f"torch=={torch.__version__}",
        "torchvision",
        "mlflow"
    ]

    # log the model
    mlflow.pytorch.log_model(
        model_to_log,
        "model",
        input_example=sample_input,
        pip_requirements=pip_requirements
    )

    # move the model back to the original device
    self.model = self.model.to(self.device)

    # plot and save training curves
    self.plot_training_curves()

    return self.history

def plot_training_curves(self):
    """
    plot and save the training curves
    """
    if len(self.history['train_loss']) == 0:
        return

    # create plots directory
    plots_dir = os.path.join(self.checkpoint_dir, 'plots')
    os.makedirs(plots_dir, exist_ok=True)

    # plot training and validation loss
    plt.figure(figsize=(10, 6))
    epochs = range(1, len(self.history['train_loss']) + 1)

```

```

plt.plot(epochs, self.history['train_loss'],
         'b-', label='Training Loss')

if len(self.history['val_loss']) > 0:
    plt.plot(epochs, self.history['val_loss'],
             'r-', label='Validation Loss')

plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.savefig(os.path.join(plots_dir, 'loss_curves.png')) # save the plot
↳ as a png file
plt.tight_layout()
plt.show()

# plot validation accuracy if available
if len(self.history['val_accuracy']) > 0:
    plt.figure(figsize=(10, 6))
    plt.plot(epochs, self.history['val_accuracy'], 'g-')
    plt.title('Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy (%)')
    plt.grid(True)
    plt.savefig(os.path.join(plots_dir, 'accuracy_curve.png')) # save
↳ the plot as a png file
plt.tight_layout()
plt.show()

# helper class for context management when not using mlflow
class DummyContextManager:
    def __enter__(self):
        return None

    def __exit__(self, *args):
        pass

```

## 1.6 Runner

```

[6]: # set random seed for reproducibility
torch.manual_seed(42)

# check for gpu availability
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f'Using device: {device}\n')

```

```

# define transformations
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

# load mnist dataset
train_dataset = torchvision.datasets.MNIST(
    root='./data',
    train=True,
    download=True,
    transform=transform
)

test_dataset = torchvision.datasets.MNIST(
    root='./data',
    train=False,
    download=True,
    transform=transform
)

# create data loaders
train_loader = DataLoader(
    train_dataset, batch_size=64, shuffle=True, pin_memory=True)
test_loader = DataLoader(test_dataset, batch_size=64,
    shuffle=False, pin_memory=True)

# create model
input_size = 28 * 28 # MNIST images are 28x28 pixels
hidden_size = [128] # one hidden layer
output_size = 10 # 10 classes (digits 0-9)
model = Perceptron(
    input_size=input_size,
    hidden_size=hidden_size,
    output_size=output_size,
    softmax_dim=-1
)

# define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

# create a learning rate scheduler for overfitting
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer,
    mode='min',

```



```

        factor=0.1,
        patience=3,
    )

    # create our trainer
    trainer = Trainer(
        model=model,
        criterion=criterion,
        optimizer=optimizer,
        train_loader=train_loader,
        test_loader=test_loader, # using test set as validation set
        device=device,
        scheduler=scheduler,
        checkpoint_dir='./checkpoints',
        experiment_name="MNIST_Perceptron_Model"
    )

    # train the model
    history = trainer.train(
        num_epochs=10,
        log_to_mlflow=True,
        early_stopping_patience=5 # enable early stopping for overfitting
    )

    # uncomment to resume training from a checkpoint
    # don't forget to change the checkpoint path
    # checkpoint_path = './checkpoints/checkpoint_epoch_2.pt'
    # start_epoch = trainer.load_checkpoint(checkpoint_path)
    # trainer.train(num_epochs=10, start_epoch=start_epoch)

    # uncomment to use the best model for inference
    # trainer.load_checkpoint('./checkpoints/best_model.pt')

```

Using device: cuda

Starting training...

```

Epoch [1/10], Step [100/938], Loss: 0.4322
Epoch [1/10], Step [200/938], Loss: 0.3024
Epoch [1/10], Step [300/938], Loss: 0.3595
Epoch [1/10], Step [400/938], Loss: 0.2423
Epoch [1/10], Step [500/938], Loss: 0.2881
Epoch [1/10], Step [600/938], Loss: 0.3294
Epoch [1/10], Step [700/938], Loss: 0.3723
Epoch [1/10], Step [800/938], Loss: 0.2187
Epoch [1/10], Step [900/938], Loss: 0.1286
Epoch [1/10], Loss: 0.3607
Validation - Epoch [1/10], Loss: 0.2217, Accuracy: 93.28%

```

New best model saved with accuracy: 93.28%

Epoch [2/10], Step [100/938], Loss: 0.4463  
Epoch [2/10], Step [200/938], Loss: 0.1599  
Epoch [2/10], Step [300/938], Loss: 0.3721  
Epoch [2/10], Step [400/938], Loss: 0.2644  
Epoch [2/10], Step [500/938], Loss: 0.2724  
Epoch [2/10], Step [600/938], Loss: 0.2156  
Epoch [2/10], Step [700/938], Loss: 0.1613  
Epoch [2/10], Step [800/938], Loss: 0.1414  
Epoch [2/10], Step [900/938], Loss: 0.1897  
Epoch [2/10], Loss: 0.1879

Validation - Epoch [2/10], Loss: 0.1493, Accuracy: 95.53%

New best model saved with accuracy: 95.53%

Epoch [3/10], Step [100/938], Loss: 0.1784  
Epoch [3/10], Step [200/938], Loss: 0.1165  
Epoch [3/10], Step [300/938], Loss: 0.1515  
Epoch [3/10], Step [400/938], Loss: 0.2542  
Epoch [3/10], Step [500/938], Loss: 0.0659  
Epoch [3/10], Step [600/938], Loss: 0.1272  
Epoch [3/10], Step [700/938], Loss: 0.3140  
Epoch [3/10], Step [800/938], Loss: 0.1646  
Epoch [3/10], Step [900/938], Loss: 0.0224  
Epoch [3/10], Loss: 0.1383

Validation - Epoch [3/10], Loss: 0.1256, Accuracy: 96.14%

New best model saved with accuracy: 96.14%

Epoch [4/10], Step [100/938], Loss: 0.3256  
Epoch [4/10], Step [200/938], Loss: 0.0810  
Epoch [4/10], Step [300/938], Loss: 0.0935  
Epoch [4/10], Step [400/938], Loss: 0.1180  
Epoch [4/10], Step [500/938], Loss: 0.1764  
Epoch [4/10], Step [600/938], Loss: 0.2124  
Epoch [4/10], Step [700/938], Loss: 0.0451  
Epoch [4/10], Step [800/938], Loss: 0.1489  
Epoch [4/10], Step [900/938], Loss: 0.0503  
Epoch [4/10], Loss: 0.1098

Validation - Epoch [4/10], Loss: 0.1031, Accuracy: 96.82%

New best model saved with accuracy: 96.82%

Epoch [5/10], Step [100/938], Loss: 0.0236  
Epoch [5/10], Step [200/938], Loss: 0.1259  
Epoch [5/10], Step [300/938], Loss: 0.0292  
Epoch [5/10], Step [400/938], Loss: 0.1083  
Epoch [5/10], Step [500/938], Loss: 0.0232  
Epoch [5/10], Step [600/938], Loss: 0.0444  
Epoch [5/10], Step [700/938], Loss: 0.0331

Epoch [5/10], Step [800/938], Loss: 0.0711  
Epoch [5/10], Step [900/938], Loss: 0.1074  
Epoch [5/10], Loss: 0.0960  
EarlyStopping counter: 1 out of 5  
Validation - Epoch [5/10], Loss: 0.1063, Accuracy: 96.51%

Epoch [6/10], Step [100/938], Loss: 0.0274  
Epoch [6/10], Step [200/938], Loss: 0.0241  
Epoch [6/10], Step [300/938], Loss: 0.0546  
Epoch [6/10], Step [400/938], Loss: 0.1778  
Epoch [6/10], Step [500/938], Loss: 0.0370  
Epoch [6/10], Step [600/938], Loss: 0.0572  
Epoch [6/10], Step [700/938], Loss: 0.0574  
Epoch [6/10], Step [800/938], Loss: 0.1473  
Epoch [6/10], Step [900/938], Loss: 0.0581  
Epoch [6/10], Loss: 0.0862  
Validation - Epoch [6/10], Loss: 0.0915, Accuracy: 97.08%  
New best model saved with accuracy: 97.08%

Epoch [7/10], Step [100/938], Loss: 0.1411  
Epoch [7/10], Step [200/938], Loss: 0.0363  
Epoch [7/10], Step [300/938], Loss: 0.0203  
Epoch [7/10], Step [400/938], Loss: 0.1467  
Epoch [7/10], Step [500/938], Loss: 0.0771  
Epoch [7/10], Step [600/938], Loss: 0.0529  
Epoch [7/10], Step [700/938], Loss: 0.0938  
Epoch [7/10], Step [800/938], Loss: 0.0370  
Epoch [7/10], Step [900/938], Loss: 0.0660  
Epoch [7/10], Loss: 0.0731  
EarlyStopping counter: 1 out of 5  
Validation - Epoch [7/10], Loss: 0.0983, Accuracy: 96.92%

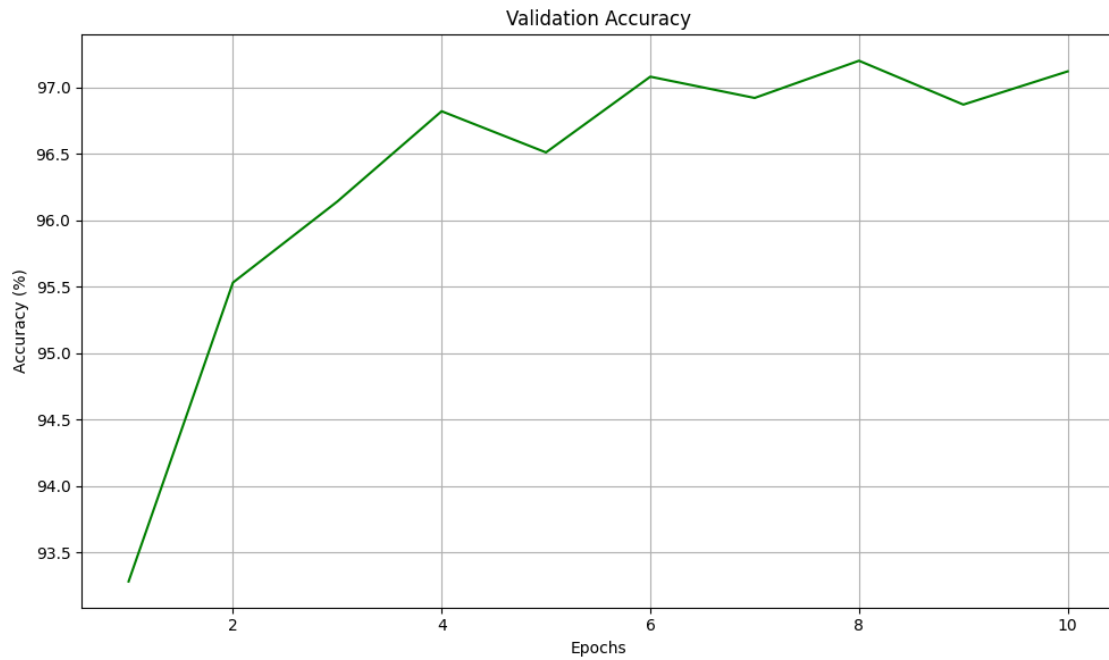
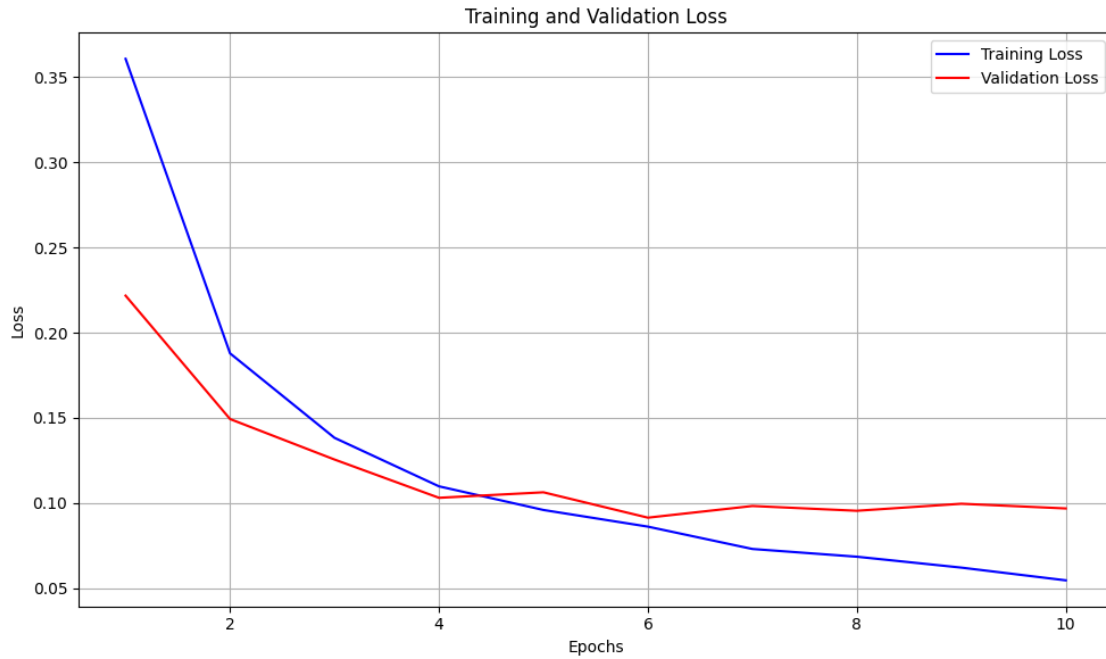
Epoch [8/10], Step [100/938], Loss: 0.2122  
Epoch [8/10], Step [200/938], Loss: 0.0411  
Epoch [8/10], Step [300/938], Loss: 0.0945  
Epoch [8/10], Step [400/938], Loss: 0.0278  
Epoch [8/10], Step [500/938], Loss: 0.0855  
Epoch [8/10], Step [600/938], Loss: 0.0477  
Epoch [8/10], Step [700/938], Loss: 0.0758  
Epoch [8/10], Step [800/938], Loss: 0.0171  
Epoch [8/10], Step [900/938], Loss: 0.1193  
Epoch [8/10], Loss: 0.0685  
EarlyStopping counter: 2 out of 5  
Validation - Epoch [8/10], Loss: 0.0955, Accuracy: 97.20%  
New best model saved with accuracy: 97.20%

Epoch [9/10], Step [100/938], Loss: 0.0374  
Epoch [9/10], Step [200/938], Loss: 0.0376

Epoch [9/10], Step [300/938], Loss: 0.0417  
Epoch [9/10], Step [400/938], Loss: 0.0091  
Epoch [9/10], Step [500/938], Loss: 0.0976  
Epoch [9/10], Step [600/938], Loss: 0.0440  
Epoch [9/10], Step [700/938], Loss: 0.0942  
Epoch [9/10], Step [800/938], Loss: 0.0112  
Epoch [9/10], Step [900/938], Loss: 0.1761  
Epoch [9/10], Loss: 0.0621  
EarlyStopping counter: 3 out of 5  
Validation - Epoch [9/10], Loss: 0.0996, Accuracy: 96.87%

Epoch [10/10], Step [100/938], Loss: 0.0189  
Epoch [10/10], Step [200/938], Loss: 0.0427  
Epoch [10/10], Step [300/938], Loss: 0.0367  
Epoch [10/10], Step [400/938], Loss: 0.0460  
Epoch [10/10], Step [500/938], Loss: 0.1362  
Epoch [10/10], Step [600/938], Loss: 0.0072  
Epoch [10/10], Step [700/938], Loss: 0.0275  
Epoch [10/10], Step [800/938], Loss: 0.0069  
Epoch [10/10], Step [900/938], Loss: 0.0329  
Epoch [10/10], Loss: 0.0547  
EarlyStopping counter: 4 out of 5  
Validation - Epoch [10/10], Loss: 0.0969, Accuracy: 97.12%  
Learning rate changed from 0.001000 to 0.000100  
Training complete!

Final Evaluation:  
Final Test Loss: 0.0969  
Final Test Accuracy: 97.12%  
Final model saved as ./checkpoints\final\_model\_20250424\_215421.pt



### 1.6.1 Other version management tools

In this part of the assignment I used MLFlow for version management but there are other tools we can use for version management. \* Weights and Biases (W&B) - Cloud-based platform for

experiment tracking, visualization and collaboration. \* Comet.ml - Cloud-based like W&B. Tracks experiments, code, hyperparameters, metrics, etc. \* Data Version Control (DVC) - Works like Git for data and model versioning \* TensorBoard - A simpler option for TensorFlow and PyTorch. Great for logging and visualization.