

Assignment-1-Part-2-4

April 25, 2025

1 Assignment 1 Part 2.4

1.1 Import necessary modules

```
[1]: import torch
import torch.nn as nn
import os
import mlflow
from datetime import datetime
import matplotlib.pyplot as plt
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import time
```

1.2 CNN Layer class

I created a CNN class with one layer

```
[2]: class CNNLayer(nn.Module):
    def __init__(self, input_channels=1, num_classes=10):
        """
        Initialize the CNNLayer class

        Parameters:
            input_channels (int): number of channels in the input image
            num_classes (int): number of classes in the output
        """
        super(CNNLayer, self).__init__()

        # define the convolutional layers
        self.conv_layer = nn.Sequential(
            # create the a convolutional layer
            nn.Conv2d(in_channels=input_channels, out_channels=32,
            ↪kernel_size=3, padding=1),
            # add a ReLU activation function which is recommended for CNNs
            nn.ReLU(),
```

```

        # add a max pooling layer which reduces the size of the image by
↪half
        nn.MaxPool2d(kernel_size=2, stride=2)
    )

    # calculate the size of the flattened layer
    # after the convolutional layer, the image is reduced to 32x14x14
    flattened_size = 32 * 14 * 14

    # define the fully connected layer
    self.fc_layer = nn.Sequential(
        nn.Flatten(), # flatten the image from 3d to 1d
        nn.Linear(flattened_size, 128),
        nn.ReLU(),
        nn.Linear(128, num_classes)
    )

    def forward(self, inputs):
        """
        Forward pass of the CNNLayer

        Parameters:
            inputs (torch.Tensor): the input image
        """

        # pass the inputs through the convolutional layer
        outputs = self.conv_layer(inputs)

        # pass the outputs through the fully connected layer
        outputs = self.fc_layer(outputs)

        return outputs

```

1.3 Trainer class: Modified

I made some small modifications so it can handle both FNN and CNN model

```

[3]: class Trainer:
    def __init__(self,
        model,
        criterion,
        optimizer,
        train_loader,
        val_loader=None,
        test_loader=None,
        device=None,
        scheduler=None,

```

```

        checkpoint_dir='./checkpoints',
        experiment_name=None):
    """
    Initialize the trainer with model, criterion, optimizer, and data_
    loaders

    Parameters:
        model: nn.Module
            The model to train
        criterion: loss function
            The loss function to use
        optimizer: torch.optim
            The optimizer to use
        train_loader: DataLoader
            The data loader for training data
        val_loader: DataLoader
            The data loader for validation data
        test_loader: DataLoader
            The data loader for test data
        device: torch.device
            The device to use for training
        scheduler: torch.optim.lr_scheduler
            Learning rate scheduler (optional)
        checkpoint_dir: str
            Directory to save checkpoints
        experiment_name: str
            Name of the experiment for MLflow tracking
    """
    self.model = model
    self.criterion = criterion
    self.optimizer = optimizer
    self.train_loader = train_loader
    self.val_loader = val_loader if val_loader is not None else test_loader
    self.test_loader = test_loader
    self.device = device if device is not None else torch.device(
        'cuda' if torch.cuda.is_available() else 'cpu')
    self.scheduler = scheduler
    self.checkpoint_dir = checkpoint_dir
    self.experiment_name = experiment_name
    self.input_size = None # will be set during first forward pass
    self.is_cnn = isinstance(self.model, CNNLayer)

    # create checkpoint directory if it doesn't exist
    os.makedirs(checkpoint_dir, exist_ok=True)

    # move model to device
    self.model.to(self.device)

```

```

    # initialize training metrics
    self.best_accuracy = 0.0
    self.best_loss = float('inf')
    self.history = {
        'train_loss': [],
        'val_loss': [],
        'val_accuracy': [],
        'learning_rates': [],
        'epoch_times': [] # add tracking of per-epoch times
    }

    # initialize performance tracking
    self.total_training_time = 0
    self.epoch_start_time = 0
    self.total_epochs_completed = 0

def train_epoch(self, epoch, num_epochs):
    """
    Train the model for one epoch
    """
    # start timing the epoch
    self.epoch_start_time = time.time()

    self.model.train()
    running_loss = 0.0

    for i, (inputs, labels) in enumerate(self.train_loader):
        # get input size from first batch if not set
        if self.input_size is None and hasattr(inputs, 'view'):
            self.input_size = inputs.view(inputs.size(0), -1).size(1)

        # prepare inputs
        if not self.is_cnn and hasattr(inputs, 'view'):
            # for image data, flatten if needed
            inputs = inputs.view(inputs.size(0), -1).to(self.device)
        else:
            inputs = inputs.to(self.device)

        labels = labels.to(self.device)

        # forward pass
        outputs = self.model(inputs)
        loss = self.criterion(outputs, labels)

        # backward and optimize
        self.optimizer.zero_grad()

```

```

        loss.backward()
        self.optimizer.step()

        # update statistics
        running_loss += loss.detach().item()

        # print progress
        if (i+1) % 100 == 0:
            print(
                f'Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/{len(self.
↪train_loader)}], Loss: {loss.detach().item():.4f}')

        # calculate average loss for the epoch
        epoch_loss = running_loss / len(self.train_loader)
        self.history['train_loss'].append(epoch_loss)

        # track learning rate
        current_lr = self.optimizer.param_groups[0]['lr']
        self.history['learning_rates'].append(current_lr)

        # calculate and store epoch time
        epoch_time = time.time() - self.epoch_start_time
        self.history['epoch_times'].append(epoch_time)
        self.total_training_time += epoch_time
        self.total_epochs_completed += 1

        # Print epoch time
        print(f'Epoch time: {epoch_time:.2f} seconds')

    return epoch_loss

def validate(self, epoch=None):
    """
    Validate the model on validation data
    """
    if self.val_loader is None:
        return None, None

    self.model.eval()
    running_loss = 0.0
    correct = 0
    total = 0

    with torch.no_grad():
        for inputs, labels in self.val_loader:
            # prepare inputs
            if not self.is_cnn and hasattr(inputs, 'view'):

```

```

        # for image data
        inputs = inputs.view(inputs.size(0), -1).to(self.device)
    else:
        inputs = inputs.to(self.device)

    labels = labels.to(self.device)

    # forward pass
    outputs = self.model(inputs)
    loss = self.criterion(outputs, labels)

    # update statistics
    running_loss += loss.item()

    # calculate accuracy
    _, predicted = torch.max(outputs.data, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

# calculate average loss and accuracy
val_loss = running_loss / len(self.val_loader)
val_accuracy = 100 * correct / total

if epoch is not None:
    self.history['val_loss'].append(val_loss)
    self.history['val_accuracy'].append(val_accuracy)

return val_loss, val_accuracy

def test(self):
    """
    Test the model on test data
    """
    if self.test_loader is None:
        return None, None

    return self.validate() # reuse validation code for testing

def save_checkpoint(self, epoch, train_loss, val_loss=None,
    ↪ val_accuracy=None, is_best=False):
    """
    Save a checkpoint of the model
    """
    checkpoint = {
        'epoch': epoch + 1,
        'model_state_dict': self.model.state_dict(),
        'optimizer_state_dict': self.optimizer.state_dict(),

```

```

        'loss': train_loss
    }

    if val_loss is not None:
        checkpoint['val_loss'] = val_loss

    if val_accuracy is not None:
        checkpoint['val_accuracy'] = val_accuracy

    if self.scheduler is not None:
        checkpoint['scheduler_state_dict'] = self.scheduler.state_dict()

    # save regular checkpoint
    checkpoint_path = os.path.join(
        self.checkpoint_dir, f'checkpoint_epoch_{epoch+1}.pt')
    torch.save(checkpoint, checkpoint_path)

    # save as best model if applicable
    if is_best:
        best_model_path = os.path.join(
            self.checkpoint_dir, 'best_model.pt')
        torch.save(checkpoint, best_model_path)
        print(f'New best model saved with accuracy: {val_accuracy:.2f}%')

    return checkpoint_path

def load_checkpoint(self, checkpoint_path):
    """
    Load a checkpoint
    """
    print(f"Loading checkpoint from {checkpoint_path}")
    checkpoint = torch.load(checkpoint_path, map_location=self.device)

    self.model.load_state_dict(checkpoint['model_state_dict'])
    self.optimizer.load_state_dict(checkpoint['optimizer_state_dict'])

    if 'scheduler_state_dict' in checkpoint and self.scheduler is not None:
        self.scheduler.load_state_dict(checkpoint['scheduler_state_dict'])

    start_epoch = checkpoint['epoch']
    loss = checkpoint.get('loss', 0)
    val_loss = checkpoint.get('val_loss', 0)
    val_accuracy = checkpoint.get('val_accuracy', 0)

    print(
        f"Checkpoint loaded - Epoch: {start_epoch}, Loss: {loss:.4f}, Val_
↵Loss: {val_loss:.4f}, Val Accuracy: {val_accuracy:.2f}%"
    )

```

```

        return start_epoch

    def train(self, num_epochs, start_epoch=0, log_to_mlflow=True,
↳early_stopping_patience=None):
        """
        Train the model for multiple epochs

        Parameters:
            num_epochs: int
                Number of epochs to train for
            start_epoch: int
                Starting epoch (useful when resuming training)
            log_to_mlflow: bool
                Whether to log metrics to MLflow
            early_stopping_patience: int
                Number of epochs to wait for improvement before stopping
        """
        # start timer for total training run
        total_run_start_time = time.time()

        # reset performance tracking for a new training run
        if start_epoch == 0:
            self.total_training_time = 0
            self.total_epochs_completed = 0

        # set mlflow experiment if provided
        if log_to_mlflow and self.experiment_name:
            mlflow.set_experiment(self.experiment_name)

        # initialize early stopping variables
        if early_stopping_patience is not None:
            early_stopping_counter = 0
            best_val_loss = float('inf')

        # start mlflow run if applicable
        run_context = mlflow.start_run(
            run_name=f"{self.model.__class__.__name__}_{datetime.now().
↳strftime('%Y%m%d_%H%M%S')}") if log_to_mlflow else DummyContextManager()

        # store previous learning rate to detect changes
        prev_lr = self.optimizer.param_groups[0]['lr']

        with run_context:
            # log parameters if using mlflow
            if log_to_mlflow:
                params = {

```



```

        "optimizer": self.optimizer.__class__.__name__,
        "learning_rate": self.optimizer.param_groups[0]['lr'],
        "num_epochs": num_epochs
    }
    # add model architecture params if available
    if hasattr(self.model, 'input_size'):
        params["input_size"] = self.model.input_size
    mlflow.log_params(params)

    print("Starting training...")
    for epoch in range(start_epoch, num_epochs):
        # train for one epoch
        print('')
        train_loss = self.train_epoch(epoch, num_epochs)
        print(
            f'Epoch [{epoch+1}/{num_epochs}], Loss: {train_loss:.4f}')

        # validate
        val_loss, val_accuracy = self.validate(epoch)

        # early stopping check for overfitting
        if early_stopping_patience is not None and val_loss is not None:
            if val_loss < best_val_loss:
                best_val_loss = val_loss
                early_stopping_counter = 0
            else:
                early_stopping_counter += 1
                print(
                    f'EarlyStopping counter: {early_stopping_counter}␣
↳out of {early_stopping_patience}')

                if early_stopping_counter >= early_stopping_patience:
                    print(
                        f'Early stopping triggered after epoch␣
↳{epoch+1}')
                    break

        # log metrics
        if log_to_mlflow:
            mlflow.log_metric("train_loss", train_loss, step=epoch)
            mlflow.log_metric("epoch_time", self.
↳history['epoch_times'][-1], step=epoch)
            if val_loss is not None:
                mlflow.log_metric("val_loss", val_loss, step=epoch)
            if val_accuracy is not None:
                mlflow.log_metric(
                    "val_accuracy", val_accuracy, step=epoch)

```

```

        # print validation results
        if val_loss is not None and val_accuracy is not None:
            print(
                f'Validation - Epoch [{epoch+1}/{num_epochs}], Loss:␣
↪{val_loss:.4f}, Accuracy: {val_accuracy:.2f}%'
            )

        # check if this is the best model so far
        is_best = False
        if val_accuracy is not None and val_accuracy > self.
↪best_accuracy:
            self.best_accuracy = val_accuracy
            is_best = True

        # save checkpoint
        self.save_checkpoint(
            epoch, train_loss, val_loss, val_accuracy, is_best)

        # update learning rate if scheduler is provided
        if self.scheduler is not None:
            if isinstance(self.scheduler, torch.optim.lr_scheduler.
↪ReduceLROnPlateau):
                self.scheduler.step(val_loss)
            else:
                self.scheduler.step()

        # check if learning rate changed
        current_lr = self.optimizer.param_groups[0]['lr']
        if current_lr != prev_lr:
            print(
                f"Learning rate changed from {prev_lr:.6f} to␣
↪{current_lr:.6f}")
            prev_lr = current_lr

        # Calculate total run time
        total_run_time = time.time() - total_run_start_time

        print("Training complete!")

        # Print performance summary
        print("\nTraining Performance Summary:")
        print(f"Total training time: {total_run_time:.2f} seconds␣
↪({total_run_time/60:.2f} minutes)")
        print(f"Total epochs completed: {self.total_epochs_completed}")
        print(f"Average time per epoch: {self.total_training_time/self.
↪total_epochs_completed:.2f} seconds")

```

```

        print(f"Fastest epoch: {min(self.history['epoch_times']):.2f}␣
↪seconds")
        print(f"Slowest epoch: {max(self.history['epoch_times']):.2f}␣
↪seconds")

        # Log performance metrics to MLflow
        if log_to_mlflow:
            mlflow.log_metric("total_training_time", total_run_time)
            mlflow.log_metric("avg_epoch_time", self.total_training_time/
↪self.total_epochs_completed)
            mlflow.log_metric("fastest_epoch_time", min(self.
↪history['epoch_times']))
            mlflow.log_metric("slowest_epoch_time", max(self.
↪history['epoch_times']))

        # final evaluation
        print("\nFinal Evaluation:")
        test_loss, test_accuracy = self.test()
        if test_loss is not None and test_accuracy is not None:
            print(f'Final Test Loss: {test_loss:.4f}')
            print(f'Final Test Accuracy: {test_accuracy:.2f}%')

        # log final metrics
        if log_to_mlflow:
            mlflow.log_metric("test_loss", test_loss)
            mlflow.log_metric("test_accuracy", test_accuracy)

        # save the final model
        model_timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
        final_model_path = os.path.join(
            self.checkpoint_dir, f"final_model_{model_timestamp}.pt")
        torch.save(self.model.state_dict(), final_model_path)
        print(f"Final model saved as {final_model_path}")

        # log the model with mlflow
        if log_to_mlflow:

            model_to_log = self.model.to('cpu')

            sample_batch = next(iter(self.train_loader))[0][:1]

            if self.is_cnn:
                # for cnn, use the first batch of the training data
                sample_input = sample_batch.numpy()
            else:
                # for fnn, use the first batch of the training data

```

```

        sample_input = sample_batch.view(sample_batch.size(0), -1).
        ↪numpy()

        # Create custom pip requirements
        pip_requirements = [
            f"torch=={torch.__version__}",
            "torchvision",
            "mlflow"
        ]

        # log the model
        mlflow.pytorch.log_model(
            model_to_log,
            "model",
            input_example=sample_input,
            pip_requirements=pip_requirements
        )

        # Move the model back to the original device
        self.model = self.model.to(self.device)

        # plot and save training curves
        self.plot_training_curves()

        return self.history

def plot_training_curves(self):
    """
    plot and save the training curves
    """
    if len(self.history['train_loss']) == 0:
        return

    # create plots directory
    plots_dir = os.path.join(self.checkpoint_dir, 'plots')
    os.makedirs(plots_dir, exist_ok=True)

    # plot training and validation loss
    plt.figure(figsize=(10, 6))
    epochs = range(1, len(self.history['train_loss']) + 1)
    plt.plot(epochs, self.history['train_loss'],
             'b-', label='Training Loss')

    if len(self.history['val_loss']) > 0:
        plt.plot(epochs, self.history['val_loss'],
             'r-', label='Validation Loss')

```

```

plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.savefig(os.path.join(plots_dir, 'loss_curves.png')) # save the plot
↳ as a png file
plt.tight_layout()
plt.show()

# plot validation accuracy if available
if len(self.history['val_accuracy']) > 0:
    plt.figure(figsize=(10, 6))
    plt.plot(epochs, self.history['val_accuracy'], 'g-')
    plt.title('Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy (%)')
    plt.grid(True)
    plt.savefig(os.path.join(plots_dir, 'accuracy_curve.png')) # save
↳ the plot as a png file
    plt.tight_layout()
    plt.show()

# plot epoch times
if len(self.history['epoch_times']) > 0:
    plt.figure(figsize=(10, 6))
    plt.plot(epochs, self.history['epoch_times'], 'm-')
    plt.title('Epoch Training Times')
    plt.xlabel('Epoch')
    plt.ylabel('Time (seconds)')
    plt.grid(True)
    plt.savefig(os.path.join(plots_dir, 'epoch_times.png')) # save the
↳ plot as a png file
    plt.tight_layout()
    plt.show()

# helper class for context management when not using mlflow
class DummyContextManager:
    def __enter__(self):
        return None

    def __exit__(self, *args):
        pass

```

1.4 Runner: Modified

I commented out the creation of FNN model and created a CNN model instead

```
[4]: # set random seed for reproducibility
torch.manual_seed(42)

# check for gpu availability
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f'Using device: {device}\n')

# define transformations for training set with augmentation
train_transform = transforms.Compose([
    transforms.RandomRotation(10),
    transforms.RandomAffine(degrees=0, translate=(0.1, 0.1), scale=(0.9, 1.1),
↪shear=5),
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

# define transformations for test set without augmentation
test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

# load mnist dataset
train_dataset = torchvision.datasets.MNIST(
    root='./data',
    train=True,
    download=True,
    transform=train_transform
)

test_dataset = torchvision.datasets.MNIST(
    root='./data',
    train=False,
    download=True,
    transform=test_transform
)

# create data loaders
train_loader = DataLoader(
    train_dataset, batch_size=64, shuffle=True, pin_memory=True)
test_loader = DataLoader(test_dataset, batch_size=64,
    shuffle=False, pin_memory=True)

# create FNN model
```

```

# input_size = 28 * 28 # MNIST images are 28x28 pixels
# hidden_size = [128] # one hidden layer
# output_size = 10 # 10 classes (digits 0-9)
# model = Perceptron(
#     input_size=input_size,
#     hidden_size=hidden_size,
#     output_size=output_size,
#     softmax_dim=-1
# )

# create CNN model
input_channels = 1
num_classes = 10
model = CNNLayer(input_channels, num_classes)

# define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

# create a learning rate scheduler for overfitting
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer,
    mode='min',
    factor=0.1,
    patience=3,
)

# create our trainer
trainer = Trainer(
    model=model,
    criterion=criterion,
    optimizer=optimizer,
    train_loader=train_loader,
    test_loader=test_loader, # using test set as validation set
    device=device,
    scheduler=scheduler,
    checkpoint_dir='./checkpoints',
    experiment_name="MNIST_CNN_Model"
)

# train the model
history = trainer.train(
    num_epochs=10,
    log_to_mlflow=True,
    early_stopping_patience=5 # enable early stopping for overfitting
)

```

```
# uncomment to resume training from a checkpoint
# don't forget to change the checkpoint path
# checkpoint_path = './checkpoints/checkpoint_epoch_2.pt'
# start_epoch = trainer.load_checkpoint(checkpoint_path)
# trainer.train(num_epochs=10, start_epoch=start_epoch)

# uncomment to use the best model for inference
# trainer.load_checkpoint('./checkpoints/best_model.pt')
```

Using device: cuda

Starting training...

```
Epoch [1/10], Step [100/938], Loss: 1.1281
Epoch [1/10], Step [200/938], Loss: 0.8483
Epoch [1/10], Step [300/938], Loss: 0.6369
Epoch [1/10], Step [400/938], Loss: 0.3395
Epoch [1/10], Step [500/938], Loss: 0.3198
Epoch [1/10], Step [600/938], Loss: 0.4834
Epoch [1/10], Step [700/938], Loss: 0.5666
Epoch [1/10], Step [800/938], Loss: 0.2994
Epoch [1/10], Step [900/938], Loss: 0.4328
Epoch time: 33.62 seconds
Epoch [1/10], Loss: 0.5883
Validation - Epoch [1/10], Loss: 0.1057, Accuracy: 97.04%
New best model saved with accuracy: 97.04%
```

```
Epoch [2/10], Step [100/938], Loss: 0.3207
Epoch [2/10], Step [200/938], Loss: 0.2179
Epoch [2/10], Step [300/938], Loss: 0.2367
Epoch [2/10], Step [400/938], Loss: 0.1585
Epoch [2/10], Step [500/938], Loss: 0.2432
Epoch [2/10], Step [600/938], Loss: 0.1588
Epoch [2/10], Step [700/938], Loss: 0.3075
Epoch [2/10], Step [800/938], Loss: 0.2538
Epoch [2/10], Step [900/938], Loss: 0.1746
Epoch time: 35.13 seconds
Epoch [2/10], Loss: 0.2529
Validation - Epoch [2/10], Loss: 0.0781, Accuracy: 97.68%
New best model saved with accuracy: 97.68%
```

```
Epoch [3/10], Step [100/938], Loss: 0.1443
Epoch [3/10], Step [200/938], Loss: 0.0725
Epoch [3/10], Step [300/938], Loss: 0.1948
Epoch [3/10], Step [400/938], Loss: 0.2072
Epoch [3/10], Step [500/938], Loss: 0.2880
Epoch [3/10], Step [600/938], Loss: 0.1831
Epoch [3/10], Step [700/938], Loss: 0.1295
```


Epoch [3/10], Step [800/938], Loss: 0.2508
Epoch [3/10], Step [900/938], Loss: 0.2199
Epoch time: 36.41 seconds
Epoch [3/10], Loss: 0.1943
EarlyStopping counter: 1 out of 5
Validation - Epoch [3/10], Loss: 0.0790, Accuracy: 97.60%

Epoch [4/10], Step [100/938], Loss: 0.2529
Epoch [4/10], Step [200/938], Loss: 0.0618
Epoch [4/10], Step [300/938], Loss: 0.1087
Epoch [4/10], Step [400/938], Loss: 0.0544
Epoch [4/10], Step [500/938], Loss: 0.0954
Epoch [4/10], Step [600/938], Loss: 0.1171
Epoch [4/10], Step [700/938], Loss: 0.1091
Epoch [4/10], Step [800/938], Loss: 0.2946
Epoch [4/10], Step [900/938], Loss: 0.0471
Epoch time: 35.26 seconds
Epoch [4/10], Loss: 0.1681
Validation - Epoch [4/10], Loss: 0.0680, Accuracy: 97.96%
New best model saved with accuracy: 97.96%

Epoch [5/10], Step [100/938], Loss: 0.1064
Epoch [5/10], Step [200/938], Loss: 0.2760
Epoch [5/10], Step [300/938], Loss: 0.0616
Epoch [5/10], Step [400/938], Loss: 0.1529
Epoch [5/10], Step [500/938], Loss: 0.1073
Epoch [5/10], Step [600/938], Loss: 0.0921
Epoch [5/10], Step [700/938], Loss: 0.0696
Epoch [5/10], Step [800/938], Loss: 0.0397
Epoch [5/10], Step [900/938], Loss: 0.2554
Epoch time: 35.20 seconds
Epoch [5/10], Loss: 0.1455
Validation - Epoch [5/10], Loss: 0.0617, Accuracy: 98.03%
New best model saved with accuracy: 98.03%

Epoch [6/10], Step [100/938], Loss: 0.1917
Epoch [6/10], Step [200/938], Loss: 0.0946
Epoch [6/10], Step [300/938], Loss: 0.1179
Epoch [6/10], Step [400/938], Loss: 0.0989
Epoch [6/10], Step [500/938], Loss: 0.0872
Epoch [6/10], Step [600/938], Loss: 0.1843
Epoch [6/10], Step [700/938], Loss: 0.1694
Epoch [6/10], Step [800/938], Loss: 0.1304
Epoch [6/10], Step [900/938], Loss: 0.1556
Epoch time: 35.65 seconds
Epoch [6/10], Loss: 0.1359
Validation - Epoch [6/10], Loss: 0.0558, Accuracy: 98.20%
New best model saved with accuracy: 98.20%

Epoch [7/10], Step [100/938], Loss: 0.0438
Epoch [7/10], Step [200/938], Loss: 0.0510
Epoch [7/10], Step [300/938], Loss: 0.1119
Epoch [7/10], Step [400/938], Loss: 0.1527
Epoch [7/10], Step [500/938], Loss: 0.1559
Epoch [7/10], Step [600/938], Loss: 0.0759
Epoch [7/10], Step [700/938], Loss: 0.1577
Epoch [7/10], Step [800/938], Loss: 0.1844
Epoch [7/10], Step [900/938], Loss: 0.0806
Epoch time: 35.05 seconds
Epoch [7/10], Loss: 0.1238
EarlyStopping counter: 1 out of 5
Validation - Epoch [7/10], Loss: 0.0571, Accuracy: 98.03%

Epoch [8/10], Step [100/938], Loss: 0.0947
Epoch [8/10], Step [200/938], Loss: 0.1240
Epoch [8/10], Step [300/938], Loss: 0.0935
Epoch [8/10], Step [400/938], Loss: 0.0807
Epoch [8/10], Step [500/938], Loss: 0.2190
Epoch [8/10], Step [600/938], Loss: 0.0678
Epoch [8/10], Step [700/938], Loss: 0.0508
Epoch [8/10], Step [800/938], Loss: 0.0694
Epoch [8/10], Step [900/938], Loss: 0.0873
Epoch time: 36.15 seconds
Epoch [8/10], Loss: 0.1177
Validation - Epoch [8/10], Loss: 0.0491, Accuracy: 98.54%
New best model saved with accuracy: 98.54%

Epoch [9/10], Step [100/938], Loss: 0.1470
Epoch [9/10], Step [200/938], Loss: 0.2309
Epoch [9/10], Step [300/938], Loss: 0.0494
Epoch [9/10], Step [400/938], Loss: 0.1545
Epoch [9/10], Step [500/938], Loss: 0.1070
Epoch [9/10], Step [600/938], Loss: 0.1042
Epoch [9/10], Step [700/938], Loss: 0.1187
Epoch [9/10], Step [800/938], Loss: 0.0433
Epoch [9/10], Step [900/938], Loss: 0.0265
Epoch time: 35.05 seconds
Epoch [9/10], Loss: 0.1102
EarlyStopping counter: 1 out of 5
Validation - Epoch [9/10], Loss: 0.0529, Accuracy: 98.27%

Epoch [10/10], Step [100/938], Loss: 0.1151
Epoch [10/10], Step [200/938], Loss: 0.1748
Epoch [10/10], Step [300/938], Loss: 0.1730
Epoch [10/10], Step [400/938], Loss: 0.2461
Epoch [10/10], Step [500/938], Loss: 0.1967

Epoch [10/10], Step [600/938], Loss: 0.1331
Epoch [10/10], Step [700/938], Loss: 0.0495
Epoch [10/10], Step [800/938], Loss: 0.1437
Epoch [10/10], Step [900/938], Loss: 0.0758
Epoch time: 35.09 seconds
Epoch [10/10], Loss: 0.1086
EarlyStopping counter: 2 out of 5
Validation - Epoch [10/10], Loss: 0.0526, Accuracy: 98.30%
Training complete!

Training Performance Summary:
Total training time: 381.48 seconds (6.36 minutes)
Total epochs completed: 10
Average time per epoch: 35.26 seconds
Fastest epoch: 33.62 seconds
Slowest epoch: 36.41 seconds

Final Evaluation:
Final Test Loss: 0.0526
Final Test Accuracy: 98.30%
Final model saved as ./checkpoints\final_model_20250425_203146.pt



