

Assignment-1-Part-2-5

April 25, 2025

1 Assignment 1 Part 2.5

1.1 Import necessary modules

```
[1]: import torch
import torch.nn as nn
import os
import mlflow
from datetime import datetime
import matplotlib.pyplot as plt
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import time
```

1.2 CNN Layer class: Modified

Added a second layer with input channel 32 and output channel 64

```
[2]: class CNNLayer(nn.Module):
    def __init__(self, input_channels=1, num_classes=10):
        """
        Initialize the CNNLayer class

        Parameters:
            input_channels (int): number of channels in the input image
            num_classes (int): number of classes in the output
        """
        super(CNNLayer, self).__init__()

        # define the convolutional layers
        self.conv_layer = nn.Sequential(
            # create the first convolutional layer
            nn.Conv2d(in_channels=input_channels, out_channels=32,
            ↪kernel_size=3, padding=1),
            # add a ReLU activation function which is recommended for CNNs
            nn.ReLU(),
```

```

        # add a max pooling layer which reduces the size of the image by
        ↪half
        nn.MaxPool2d(kernel_size=2, stride=2),

        # create the second convolutional layer
        nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3,
        ↪padding=1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2)
    )

    # calculate the size of the flattened layer
    # after the convolutional layer, the image is reduced to 64x7x7
    flattened_size = 64 * 7 * 7

    # define the fully connected layer
    self.fc_layer = nn.Sequential(
        nn.Flatten(), # flatten the image from 3d to 1d
        nn.Linear(flattened_size, 128),
        nn.ReLU(),
        nn.Linear(128, num_classes)
    )

    def forward(self, inputs):
        """
        Forward pass of the CNNLayer

        Parameters:
            inputs (torch.Tensor): the input image
        """

        # pass the inputs through the convolutional layer
        outputs = self.conv_layer(inputs)

        # pass the outputs through the fully connected layer
        outputs = self.fc_layer(outputs)

        return outputs

```

1.3 Trainer class

```

[3]: class Trainer:
    def __init__(self,
                 model,
                 criterion,
                 optimizer,
                 train_loader,

```

```

        val_loader=None,
        test_loader=None,
        device=None,
        scheduler=None,
        checkpoint_dir='./checkpoints',
        experiment_name=None):
    """
    Initialize the trainer with model, criterion, optimizer, and data_
    ↪loaders

    Parameters:
        model: nn.Module
            The model to train
        criterion: loss function
            The loss function to use
        optimizer: torch.optim
            The optimizer to use
        train_loader: DataLoader
            The data loader for training data
        val_loader: DataLoader
            The data loader for validation data
        test_loader: DataLoader
            The data loader for test data
        device: torch.device
            The device to use for training
        scheduler: torch.optim.lr_scheduler
            Learning rate scheduler (optional)
        checkpoint_dir: str
            Directory to save checkpoints
        experiment_name: str
            Name of the experiment for MLflow tracking
    """
    self.model = model
    self.criterion = criterion
    self.optimizer = optimizer
    self.train_loader = train_loader
    self.val_loader = val_loader if val_loader is not None else test_loader
    self.test_loader = test_loader
    self.device = device if device is not None else torch.device(
        'cuda' if torch.cuda.is_available() else 'cpu')
    self.scheduler = scheduler
    self.checkpoint_dir = checkpoint_dir
    self.experiment_name = experiment_name
    self.input_size = None # will be set during first forward pass
    self.is_cnn = isinstance(self.model, CNNLayer)

    # create checkpoint directory if it doesn't exist

```

```

os.makedirs(checkpoint_dir, exist_ok=True)

# move model to device
self.model.to(self.device)

# initialize training metrics
self.best_accuracy = 0.0
self.best_loss = float('inf')
self.history = {
    'train_loss': [],
    'val_loss': [],
    'val_accuracy': [],
    'learning_rates': [],
    'epoch_times': [] # add tracking of per-epoch times
}

# initialize performance tracking
self.total_training_time = 0
self.epoch_start_time = 0
self.total_epochs_completed = 0

def train_epoch(self, epoch, num_epochs):
    """
    Train the model for one epoch
    """
    # start timing the epoch
    self.epoch_start_time = time.time()

    self.model.train()
    running_loss = 0.0

    for i, (inputs, labels) in enumerate(self.train_loader):
        # get input size from first batch if not set
        if self.input_size is None and hasattr(inputs, 'view'):
            self.input_size = inputs.view(inputs.size(0), -1).size(1)

        # prepare inputs
        if not self.is_cnn and hasattr(inputs, 'view'):
            # for image data, flatten if needed
            inputs = inputs.view(inputs.size(0), -1).to(self.device)
        else:
            inputs = inputs.to(self.device)

        labels = labels.to(self.device)

        # forward pass
        outputs = self.model(inputs)

```

```

        loss = self.criterion(outputs, labels)

        # backward and optimize
        self.optimizer.zero_grad()
        loss.backward()
        self.optimizer.step()

        # update statistics
        running_loss += loss.detach().item()

        # print progress
        if (i+1) % 100 == 0:
            print(
                f'Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/{len(self.
→train_loader)}], Loss: {loss.detach().item():.4f}')

        # calculate average loss for the epoch
        epoch_loss = running_loss / len(self.train_loader)
        self.history['train_loss'].append(epoch_loss)

        # track learning rate
        current_lr = self.optimizer.param_groups[0]['lr']
        self.history['learning_rates'].append(current_lr)

        # calculate and store epoch time
        epoch_time = time.time() - self.epoch_start_time
        self.history['epoch_times'].append(epoch_time)
        self.total_training_time += epoch_time
        self.total_epochs_completed += 1

        # Print epoch time
        print(f'Epoch time: {epoch_time:.2f} seconds')

    return epoch_loss

def validate(self, epoch=None):
    """
    Validate the model on validation data
    """
    if self.val_loader is None:
        return None, None

    self.model.eval()
    running_loss = 0.0
    correct = 0
    total = 0

```

```

with torch.no_grad():
    for inputs, labels in self.val_loader:
        # prepare inputs
        if not self.is_cnn and hasattr(inputs, 'view'):
            # for image data
            inputs = inputs.view(inputs.size(0), -1).to(self.device)
        else:
            inputs = inputs.to(self.device)

        labels = labels.to(self.device)

        # forward pass
        outputs = self.model(inputs)
        loss = self.criterion(outputs, labels)

        # update statistics
        running_loss += loss.item()

        # calculate accuracy
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    # calculate average loss and accuracy
    val_loss = running_loss / len(self.val_loader)
    val_accuracy = 100 * correct / total

    if epoch is not None:
        self.history['val_loss'].append(val_loss)
        self.history['val_accuracy'].append(val_accuracy)

    return val_loss, val_accuracy

def test(self):
    """
    Test the model on test data
    """
    if self.test_loader is None:
        return None, None

    return self.validate() # reuse validation code for testing

def save_checkpoint(self, epoch, train_loss, val_loss=None,
    ↪ val_accuracy=None, is_best=False):
    """
    Save a checkpoint of the model
    """

```

```

checkpoint = {
    'epoch': epoch + 1,
    'model_state_dict': self.model.state_dict(),
    'optimizer_state_dict': self.optimizer.state_dict(),
    'loss': train_loss
}

if val_loss is not None:
    checkpoint['val_loss'] = val_loss

if val_accuracy is not None:
    checkpoint['val_accuracy'] = val_accuracy

if self.scheduler is not None:
    checkpoint['scheduler_state_dict'] = self.scheduler.state_dict()

# save regular checkpoint
checkpoint_path = os.path.join(
    self.checkpoint_dir, f'checkpoint_epoch_{epoch+1}.pt')
torch.save(checkpoint, checkpoint_path)

# save as best model if applicable
if is_best:
    best_model_path = os.path.join(
        self.checkpoint_dir, 'best_model.pt')
    torch.save(checkpoint, best_model_path)
    print(f'New best model saved with accuracy: {val_accuracy:.2f}%')

return checkpoint_path

def load_checkpoint(self, checkpoint_path):
    """
    Load a checkpoint
    """
    print(f"Loading checkpoint from {checkpoint_path}")
    checkpoint = torch.load(checkpoint_path, map_location=self.device)

    self.model.load_state_dict(checkpoint['model_state_dict'])
    self.optimizer.load_state_dict(checkpoint['optimizer_state_dict'])

    if 'scheduler_state_dict' in checkpoint and self.scheduler is not None:
        self.scheduler.load_state_dict(checkpoint['scheduler_state_dict'])

    start_epoch = checkpoint['epoch']
    loss = checkpoint.get('loss', 0)
    val_loss = checkpoint.get('val_loss', 0)
    val_accuracy = checkpoint.get('val_accuracy', 0)

```

```

print(
    f"Checkpoint loaded - Epoch: {start_epoch}, Loss: {loss:.4f}, Val_
↳ Loss: {val_loss:.4f}, Val Accuracy: {val_accuracy:.2f}%"

return start_epoch

def train(self, num_epochs, start_epoch=0, log_to_mlflow=True,
↳ early_stopping_patience=None):
    """
    Train the model for multiple epochs

    Parameters:
        num_epochs: int
            Number of epochs to train for
        start_epoch: int
            Starting epoch (useful when resuming training)
        log_to_mlflow: bool
            Whether to log metrics to MLflow
        early_stopping_patience: int
            Number of epochs to wait for improvement before stopping
    """
    # start timer for total training run
    total_run_start_time = time.time()

    # reset performance tracking for a new training run
    if start_epoch == 0:
        self.total_training_time = 0
        self.total_epochs_completed = 0

    # set mlflow experiment if provided
    if log_to_mlflow and self.experiment_name:
        mlflow.set_experiment(self.experiment_name)

    # initialize early stopping variables
    if early_stopping_patience is not None:
        early_stopping_counter = 0
        best_val_loss = float('inf')

    # start mlflow run if applicable
    run_context = mlflow.start_run(
        run_name=f"{self.model.__class__.__name__}_{datetime.now().
↳ strftime('%Y%m%d_%H%M%S')}") if log_to_mlflow else DummyContextManager()

    # store previous learning rate to detect changes
    prev_lr = self.optimizer.param_groups[0]['lr']

```



```

with run_context:
    # log parameters if using mlflow
    if log_to_mlflow:
        params = {
            "optimizer": self.optimizer.__class__.__name__,
            "learning_rate": self.optimizer.param_groups[0]['lr'],
            "num_epochs": num_epochs
        }
        # add model architecture params if available
        if hasattr(self.model, 'input_size'):
            params["input_size"] = self.model.input_size
        mlflow.log_params(params)

print("Starting training...")
for epoch in range(start_epoch, num_epochs):
    # train for one epoch
    print('')
    train_loss = self.train_epoch(epoch, num_epochs)
    print(
        f'Epoch [{epoch+1}/{num_epochs}], Loss: {train_loss:.4f}')

    # validate
    val_loss, val_accuracy = self.validate(epoch)

    # early stopping check for overfitting
    if early_stopping_patience is not None and val_loss is not None:
        if val_loss < best_val_loss:
            best_val_loss = val_loss
            early_stopping_counter = 0
        else:
            early_stopping_counter += 1
            print(
                f'EarlyStopping counter: {early_stopping_counter}␣
↳out of {early_stopping_patience}')

            if early_stopping_counter >= early_stopping_patience:
                print(
                    f'Early stopping triggered after epoch␣
↳{epoch+1}')
                break

    # log metrics
    if log_to_mlflow:
        mlflow.log_metric("train_loss", train_loss, step=epoch)
        mlflow.log_metric("epoch_time", self.
↳history['epoch_times'][-1], step=epoch)
        if val_loss is not None:

```

```

        mlflow.log_metric("val_loss", val_loss, step=epoch)
    if val_accuracy is not None:
        mlflow.log_metric(
            "val_accuracy", val_accuracy, step=epoch)

    # print validation results
    if val_loss is not None and val_accuracy is not None:
        print(
            f'Validation - Epoch [{epoch+1}/{num_epochs}], Loss:␣
↪{val_loss:.4f}, Accuracy: {val_accuracy:.2f}%' )

    # check if this is the best model so far
    is_best = False
    if val_accuracy is not None and val_accuracy > self.
↪best_accuracy:
        self.best_accuracy = val_accuracy
        is_best = True

    # save checkpoint
    self.save_checkpoint(
        epoch, train_loss, val_loss, val_accuracy, is_best)

    # update learning rate if scheduler is provided
    if self.scheduler is not None:
        if isinstance(self.scheduler, torch.optim.lr_scheduler.
↪ReduceLROnPlateau):
            self.scheduler.step(val_loss)
        else:
            self.scheduler.step()

    # check if learning rate changed
    current_lr = self.optimizer.param_groups[0]['lr']
    if current_lr != prev_lr:
        print(
            f"Learning rate changed from {prev_lr:.6f} to␣
↪{current_lr:.6f}")
        prev_lr = current_lr

    # Calculate total run time
    total_run_time = time.time() - total_run_start_time

    print("Training complete!")

    # Print performance summary
    print("\nTraining Performance Summary:")

```

```

        print(f"Total training time: {total_run_time:.2f} seconds_
↪({total_run_time/60:.2f} minutes)")
        print(f"Total epochs completed: {self.total_epochs_completed}")
        print(f"Average time per epoch: {self.total_training_time/self.
↪total_epochs_completed:.2f} seconds")
        print(f"Fastest epoch: {min(self.history['epoch_times']):.2f}_
↪seconds")
        print(f"Slowest epoch: {max(self.history['epoch_times']):.2f}_
↪seconds")

        # Log performance metrics to MLflow
        if log_to_mlflow:
            mlflow.log_metric("total_training_time", total_run_time)
            mlflow.log_metric("avg_epoch_time", self.total_training_time/
↪self.total_epochs_completed)
            mlflow.log_metric("fastest_epoch_time", min(self.
↪history['epoch_times']))
            mlflow.log_metric("slowest_epoch_time", max(self.
↪history['epoch_times']))

        # final evaluation
        print("\nFinal Evaluation:")
        test_loss, test_accuracy = self.test()
        if test_loss is not None and test_accuracy is not None:
            print(f'Final Test Loss: {test_loss:.4f}')
            print(f'Final Test Accuracy: {test_accuracy:.2f}%')

        # log final metrics
        if log_to_mlflow:
            mlflow.log_metric("test_loss", test_loss)
            mlflow.log_metric("test_accuracy", test_accuracy)

        # save the final model
        model_timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
        final_model_path = os.path.join(
            self.checkpoint_dir, f"final_model_{model_timestamp}.pt")
        torch.save(self.model.state_dict(), final_model_path)
        print(f"Final model saved as {final_model_path}")

        # log the model with mlflow
        if log_to_mlflow:

            model_to_log = self.model.to('cpu')

            sample_batch = next(iter(self.train_loader))[0][:1]

```

```

        if self.is_cnn:
            # for cnn, use the first batch of the training data
            sample_input = sample_batch.numpy()
        else:
            # for fnn, use the first batch of the training data
            sample_input = sample_batch.view(sample_batch.size(0), -1).
→numpy()

        # Create custom pip requirements
        pip_requirements = [
            f"torch=={torch.__version__}",
            "torchvision",
            "mlflow"
        ]

        # log the model
        mlflow.pytorch.log_model(
            model_to_log,
            "model",
            input_example=sample_input,
            pip_requirements=pip_requirements
        )

        # Move the model back to the original device
        self.model = self.model.to(self.device)

        # plot and save training curves
        self.plot_training_curves()

        return self.history

def plot_training_curves(self):
    """
    plot and save the training curves
    """
    if len(self.history['train_loss']) == 0:
        return

    # create plots directory
    plots_dir = os.path.join(self.checkpoint_dir, 'plots')
    os.makedirs(plots_dir, exist_ok=True)

    # plot training and validation loss
    plt.figure(figsize=(10, 6))
    epochs = range(1, len(self.history['train_loss']) + 1)
    plt.plot(epochs, self.history['train_loss'],
             'b-', label='Training Loss')

```

```

        if len(self.history['val_loss']) > 0:
            plt.plot(epochs, self.history['val_loss'],
                     'r-', label='Validation Loss')

    plt.title('Training and Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()
    plt.grid(True)
    plt.savefig(os.path.join(plots_dir, 'loss_curves.png')) # save the plot
    ↪ as a png file
    plt.tight_layout()
    plt.show()

    # plot validation accuracy if available
    if len(self.history['val_accuracy']) > 0:
        plt.figure(figsize=(10, 6))
        plt.plot(epochs, self.history['val_accuracy'], 'g-')
        plt.title('Validation Accuracy')
        plt.xlabel('Epochs')
        plt.ylabel('Accuracy (%)')
        plt.grid(True)
        plt.savefig(os.path.join(plots_dir, 'accuracy_curve.png')) # save
    ↪ the plot as a png file
        plt.tight_layout()
        plt.show()

    # plot epoch times
    if len(self.history['epoch_times']) > 0:
        plt.figure(figsize=(10, 6))
        plt.plot(epochs, self.history['epoch_times'], 'm-')
        plt.title('Epoch Training Times')
        plt.xlabel('Epoch')
        plt.ylabel('Time (seconds)')
        plt.grid(True)
        plt.savefig(os.path.join(plots_dir, 'epoch_times.png')) # save the
    ↪ plot as a png file
        plt.tight_layout()
        plt.show()

# helper class for context management when not using mlflow
class DummyContextManager:
    def __enter__(self):
        return None

```

```
def __exit__(self, *args):  
    pass
```

1.4 Runner

```
[4]: # set random seed for reproducibility  
torch.manual_seed(42)  
  
# check for gpu availability  
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
print(f'Using device: {device}\n')  
  
# define transformations for training set with augmentation  
train_transform = transforms.Compose([  
    transforms.RandomRotation(10),  
    transforms.RandomAffine(degrees=0, translate=(0.1, 0.1), scale=(0.9, 1.1),  
↪shear=5),  
    transforms.ToTensor(),  
    transforms.Normalize((0.5,), (0.5,))  
)  
  
# define transformations for test set without augmentation  
test_transform = transforms.Compose([  
    transforms.ToTensor(),  
    transforms.Normalize((0.5,), (0.5,))  
)  
  
# load mnist dataset  
train_dataset = torchvision.datasets.MNIST(  
    root='./data',  
    train=True,  
    download=True,  
    transform=train_transform  
)  
  
test_dataset = torchvision.datasets.MNIST(  
    root='./data',  
    train=False,  
    download=True,  
    transform=test_transform  
)  
  
# create data loaders  
train_loader = DataLoader(  
    train_dataset, batch_size=64, shuffle=True, pin_memory=True)  
test_loader = DataLoader(test_dataset, batch_size=64,  
    shuffle=False, pin_memory=True)
```

```

# create FNN model
# input_size = 28 * 28 # MNIST images are 28x28 pixels
# hidden_size = [128] # one hidden layer
# output_size = 10 # 10 classes (digits 0-9)
# model = Perceptron(
#     input_size=input_size,
#     hidden_size=hidden_size,
#     output_size=output_size,
#     softmax_dim=-1
# )

# create CNN model
input_channels = 1
num_classes = 10
model = CNNLayer(input_channels, num_classes)

# define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

# create a learning rate scheduler for overfitting
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer,
    mode='min',
    factor=0.1,
    patience=3,
)

# create our trainer
trainer = Trainer(
    model=model,
    criterion=criterion,
    optimizer=optimizer,
    train_loader=train_loader,
    test_loader=test_loader, # using test set as validation set
    device=device,
    scheduler=scheduler,
    checkpoint_dir='./checkpoints',
    experiment_name="MNIST_CNN_Model"
)

# train the model
history = trainer.train(
    num_epochs=10,
    log_to_mlflow=True,
    early_stopping_patience=5 # enable early stopping for overfitting

```

```
)

# uncomment to resume training from a checkpoint
# don't forget to change the checkpoint path
# checkpoint_path = './checkpoints/checkpoint_epoch_2.pt'
# start_epoch = trainer.load_checkpoint(checkpoint_path)
# trainer.train(num_epochs=10, start_epoch=start_epoch)

# uncomment to use the best model for inference
# trainer.load_checkpoint('./checkpoints/best_model.pt')
```

Using device: cuda

Starting training...

```
Epoch [1/10], Step [100/938], Loss: 0.5428
Epoch [1/10], Step [200/938], Loss: 0.2344
Epoch [1/10], Step [300/938], Loss: 0.2716
Epoch [1/10], Step [400/938], Loss: 0.1592
Epoch [1/10], Step [500/938], Loss: 0.0947
Epoch [1/10], Step [600/938], Loss: 0.1205
Epoch [1/10], Step [700/938], Loss: 0.1551
Epoch [1/10], Step [800/938], Loss: 0.2663
Epoch [1/10], Step [900/938], Loss: 0.1149
Epoch time: 33.48 seconds
Epoch [1/10], Loss: 0.3235
Validation - Epoch [1/10], Loss: 0.0802, Accuracy: 97.38%
New best model saved with accuracy: 97.38%
```

```
Epoch [2/10], Step [100/938], Loss: 0.1157
Epoch [2/10], Step [200/938], Loss: 0.1394
Epoch [2/10], Step [300/938], Loss: 0.0645
Epoch [2/10], Step [400/938], Loss: 0.1535
Epoch [2/10], Step [500/938], Loss: 0.1854
Epoch [2/10], Step [600/938], Loss: 0.1136
Epoch [2/10], Step [700/938], Loss: 0.0622
Epoch [2/10], Step [800/938], Loss: 0.0549
Epoch [2/10], Step [900/938], Loss: 0.0481
Epoch time: 33.06 seconds
Epoch [2/10], Loss: 0.1097
Validation - Epoch [2/10], Loss: 0.0444, Accuracy: 98.55%
New best model saved with accuracy: 98.55%
```

```
Epoch [3/10], Step [100/938], Loss: 0.0816
Epoch [3/10], Step [200/938], Loss: 0.1400
Epoch [3/10], Step [300/938], Loss: 0.0290
Epoch [3/10], Step [400/938], Loss: 0.1129
Epoch [3/10], Step [500/938], Loss: 0.0771
```


Epoch [3/10], Step [600/938], Loss: 0.0300
Epoch [3/10], Step [700/938], Loss: 0.0950
Epoch [3/10], Step [800/938], Loss: 0.0235
Epoch [3/10], Step [900/938], Loss: 0.0194
Epoch time: 33.34 seconds
Epoch [3/10], Loss: 0.0838
Validation - Epoch [3/10], Loss: 0.0338, Accuracy: 98.84%
New best model saved with accuracy: 98.84%

Epoch [4/10], Step [100/938], Loss: 0.0925
Epoch [4/10], Step [200/938], Loss: 0.0752
Epoch [4/10], Step [300/938], Loss: 0.0717
Epoch [4/10], Step [400/938], Loss: 0.0456
Epoch [4/10], Step [500/938], Loss: 0.0360
Epoch [4/10], Step [600/938], Loss: 0.0411
Epoch [4/10], Step [700/938], Loss: 0.1444
Epoch [4/10], Step [800/938], Loss: 0.4482
Epoch [4/10], Step [900/938], Loss: 0.0624
Epoch time: 33.92 seconds
Epoch [4/10], Loss: 0.0675
Validation - Epoch [4/10], Loss: 0.0318, Accuracy: 98.88%
New best model saved with accuracy: 98.88%

Epoch [5/10], Step [100/938], Loss: 0.0836
Epoch [5/10], Step [200/938], Loss: 0.0077
Epoch [5/10], Step [300/938], Loss: 0.0193
Epoch [5/10], Step [400/938], Loss: 0.1097
Epoch [5/10], Step [500/938], Loss: 0.0423
Epoch [5/10], Step [600/938], Loss: 0.0346
Epoch [5/10], Step [700/938], Loss: 0.1023
Epoch [5/10], Step [800/938], Loss: 0.0970
Epoch [5/10], Step [900/938], Loss: 0.0697
Epoch time: 33.23 seconds
Epoch [5/10], Loss: 0.0614
Validation - Epoch [5/10], Loss: 0.0275, Accuracy: 99.06%
New best model saved with accuracy: 99.06%

Epoch [6/10], Step [100/938], Loss: 0.0731
Epoch [6/10], Step [200/938], Loss: 0.1700
Epoch [6/10], Step [300/938], Loss: 0.0205
Epoch [6/10], Step [400/938], Loss: 0.0665
Epoch [6/10], Step [500/938], Loss: 0.1214
Epoch [6/10], Step [600/938], Loss: 0.0443
Epoch [6/10], Step [700/938], Loss: 0.0572
Epoch [6/10], Step [800/938], Loss: 0.0399
Epoch [6/10], Step [900/938], Loss: 0.0768
Epoch time: 33.45 seconds
Epoch [6/10], Loss: 0.0557

EarlyStopping counter: 1 out of 5
Validation - Epoch [6/10], Loss: 0.0279, Accuracy: 99.09%
New best model saved with accuracy: 99.09%

Epoch [7/10], Step [100/938], Loss: 0.0196
Epoch [7/10], Step [200/938], Loss: 0.0110
Epoch [7/10], Step [300/938], Loss: 0.0103
Epoch [7/10], Step [400/938], Loss: 0.0212
Epoch [7/10], Step [500/938], Loss: 0.0073
Epoch [7/10], Step [600/938], Loss: 0.0175
Epoch [7/10], Step [700/938], Loss: 0.0767
Epoch [7/10], Step [800/938], Loss: 0.0050
Epoch [7/10], Step [900/938], Loss: 0.0097
Epoch time: 33.17 seconds
Epoch [7/10], Loss: 0.0514
Validation - Epoch [7/10], Loss: 0.0206, Accuracy: 99.23%
New best model saved with accuracy: 99.23%

Epoch [8/10], Step [100/938], Loss: 0.1658
Epoch [8/10], Step [200/938], Loss: 0.0222
Epoch [8/10], Step [300/938], Loss: 0.1191
Epoch [8/10], Step [400/938], Loss: 0.0175
Epoch [8/10], Step [500/938], Loss: 0.1153
Epoch [8/10], Step [600/938], Loss: 0.0200
Epoch [8/10], Step [700/938], Loss: 0.0269
Epoch [8/10], Step [800/938], Loss: 0.0572
Epoch [8/10], Step [900/938], Loss: 0.0283
Epoch time: 33.75 seconds
Epoch [8/10], Loss: 0.0485
EarlyStopping counter: 1 out of 5
Validation - Epoch [8/10], Loss: 0.0211, Accuracy: 99.26%
New best model saved with accuracy: 99.26%


Epoch [9/10], Step [100/938], Loss: 0.0122
Epoch [9/10], Step [200/938], Loss: 0.0044
Epoch [9/10], Step [300/938], Loss: 0.0395
Epoch [9/10], Step [400/938], Loss: 0.0713
Epoch [9/10], Step [500/938], Loss: 0.0421
Epoch [9/10], Step [600/938], Loss: 0.1049
Epoch [9/10], Step [700/938], Loss: 0.0322
Epoch [9/10], Step [800/938], Loss: 0.0912
Epoch [9/10], Step [900/938], Loss: 0.0563
Epoch time: 32.95 seconds
Epoch [9/10], Loss: 0.0467
EarlyStopping counter: 2 out of 5
Validation - Epoch [9/10], Loss: 0.0266, Accuracy: 99.10%


Epoch [10/10], Step [100/938], Loss: 0.0783

Epoch [10/10], Step [200/938], Loss: 0.0027
Epoch [10/10], Step [300/938], Loss: 0.0053
Epoch [10/10], Step [400/938], Loss: 0.0073
Epoch [10/10], Step [500/938], Loss: 0.0109
Epoch [10/10], Step [600/938], Loss: 0.0067
Epoch [10/10], Step [700/938], Loss: 0.0664
Epoch [10/10], Step [800/938], Loss: 0.0197
Epoch [10/10], Step [900/938], Loss: 0.0768
Epoch time: 33.09 seconds
Epoch [10/10], Loss: 0.0439
EarlyStopping counter: 3 out of 5
Validation - Epoch [10/10], Loss: 0.0229, Accuracy: 99.24%
Training complete!

Training Performance Summary:
Total training time: 364.21 seconds (6.07 minutes)
Total epochs completed: 10
Average time per epoch: 33.34 seconds
Fastest epoch: 32.95 seconds
Slowest epoch: 33.92 seconds

Final Evaluation:
Final Test Loss: 0.0229
Final Test Accuracy: 99.24%
Final model saved as ./checkpoints\final_model_20250425_210824.pt





1.4.1 How different model architecture affects the result? What type of features are typically detected by the later convolutional layers compared to the first convolutional layer?

Deeper networks (more layers) can capture more complex pattern, while wider layers (more filter per layer) can capture more variety of patterns with the same level of abstraction. In our case adding more layer can increase the performance but also can be a bit overkill because we have relatively simple task. Adding more layer (third and fourth) will only give us a small increase to performance.

The first layer detects visual features such as edges, lines, and simple texture with different orientations. The second layer uses the first layers simple features to detect more complex patterns.